

## AI Backend Engineer

Airlearn is a next-gen AI-powered language learning app that uses autonomous agents and LLMs to create personalized, immersive learning experiences. We're reimagining how people learn languages with smart conversational agents, real-time feedback, and adaptive content that actually sticks. Built by a small, fast-moving team in Bangalore, we're on a mission to make language learning intuitive, effective, and fun.

We're based out of Domlur, Bangalore and are assembling a team of makers, hackers, and builders to push the frontier of what autonomous systems can do in the real world.

### What You'll Be Doing:

Build, ship, and experiment with LLM-powered agentic systems

Design tools + memory flows + feedback loops that make our AI smart(er)

Work with LangChain, OpenAI, Anthropic APIs, vector DBs, and toolchains

Collaborate on building reusable, modular workflows for reasoning + action

Break things. Fix things. Repeat. Own what you build.

### What We're Looking For

1-2 years of experience working on cool ML/AI projects

You've built something with LLMs (GPT, Claude, Mistral, etc.) - even cooler if it was weird

Python, Go, APIs, async workflows? That's your jam

Bonus: played with RAG, LangChain, CrewAI, or built mini agents on weekends

You love startup chaos and hate unnecessary meetings

### Nice-to-Haves

Side projects that got traction (or GitHub stars 🙄)

Meme-worthy prompt engineering skills

Worked at an AI startup or shipped LLM products

A healthy obsession with what agents can (and can't yet) do

Full ownership of features and agent behavior

A team that actually loves AI - not just the hype

Office in Domlur with cold brew, whiteboards, and weird ideas